



MPI Parallelization, Part I

parallel::distributed::Triangulation

Denis Davydov (denis.davydov@fau.de)

[Luca Heltai \(luca.heltai@sissa.it\)](mailto:luca.heltai@sissa.it)

25 February - 1 March 2019




FRIEDRICH-ALEXANDER
UNIVERSITÄT
ERLANGEN-NÜRNBERG

TECHNISCHE FAKULTÄT



Parallel computing model: MPI

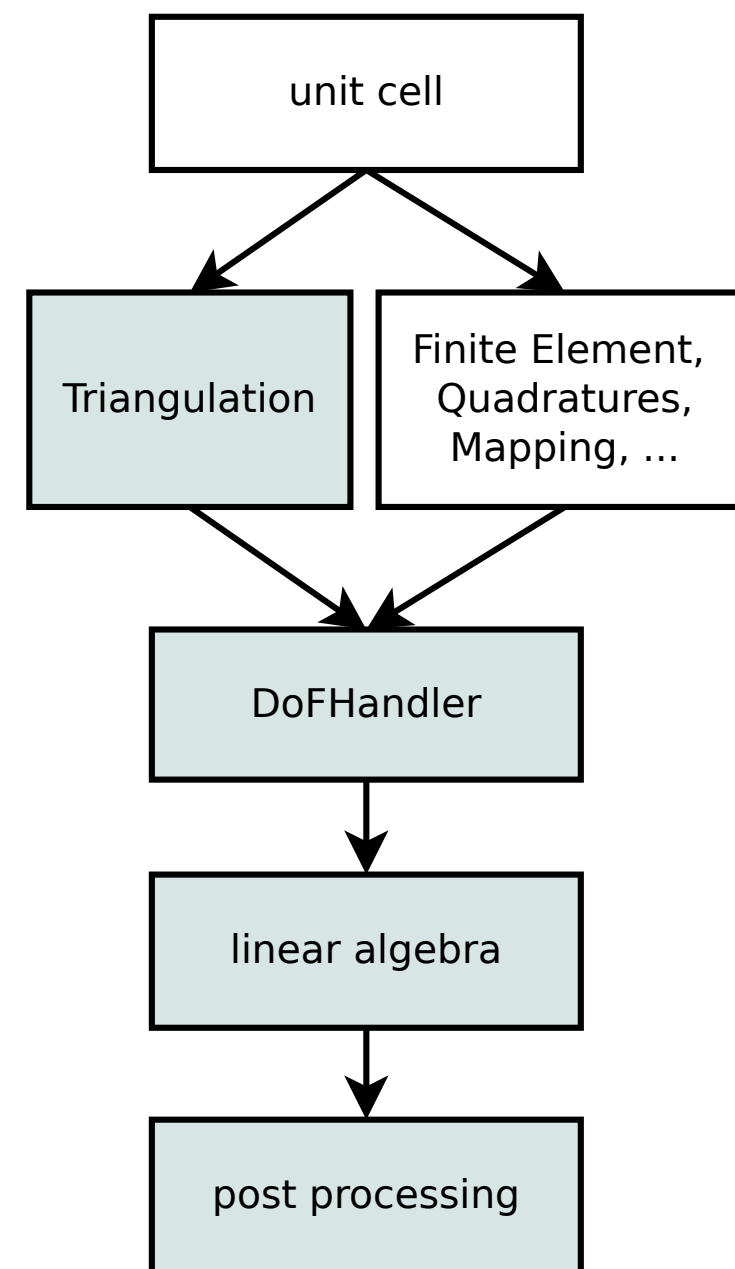


General Considerations

- Goal: get the solution faster!
- If FEM with <500.000 dofs, and 2d, use direct solver!
- If you need more, then you have to **SPLIT** the work
 - **Distributed data** storage everywhere
 - need special data structures
 - **Efficient algorithms**
 - not depending on total problem size
 - **“Localize” and “hide” communication**
 - point-to-point communication, nonblocking sends and receives

[Needs to be parallelized:


1. Triangulation (mesh with associated data)
— hard: distributed storage, new algorithms
2. DoFHandler (manages degrees of freedom)
— hard: find global numbering of DoFs
3. Linear Algebra (matrices, vectors, solvers)
— use existing library
4. Postprocessing (error estimation, solution transfer, output, ...)
— do work on local mesh, communicate



How to Parallelize?

Option 1: Domain Decomposition


- 🐾 Split up problem on PDE level
- 🐾 Solve subproblems independently
- 🐾 Converges against global solution
- 🐾 Problems:
 - 🐾 Boundary conditions are problem dependent:
 - ~> sometimes difficult!
 - ~> no black box approach!
 - 🐾 Without coarse grid solver:
 - condition number grows with $\#$ subdomains
 - ~> no linear scaling with number of CPUs!



How to Parallelize?

Option 2: Algebraic Splitting

🐾 Split up mesh between processors:



🐾 Assemble logically global linear system (distributed storage):

$$\begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix} \begin{pmatrix} \cdot \\ \cdot \\ \cdot \end{pmatrix} = \begin{pmatrix} \cdot \\ \cdot \\ \cdot \end{pmatrix}$$

🐾 Solve using iterative linear solvers in parallel

🐾 Advantages:


🐾 Looks like serial program to the user

🐾 Linear scaling possible (with good preconditioner)

Partitioning

Optimal partitioning (coloring of cells):

- 🐾 same size per region
~> even distribution of work
- 🐾 minimize interface between region
~> reduce communication






Optimal partitioning is an NP-hard
graph partitioning problem.

- 🐾 Typically done: heuristics (existing tools: METIS)
- 🐾 Problem: worse than linear runtime
- 🐾 Large graphs: several minutes, memory restrictions
- ~> Alternative: avoid graph partitioning

Partitioning with “Space filling curves”

- 🐾 *p4est* library: parallel quad-/octrees
- 🐾 Store refinement flags from a base mesh
- 🐾 Based on space-filling curves
- 🐾 Very good scalability




Burstedde, Wilcox, and Ghattas.

p4est: Scalable algorithms for parallel adaptive mesh refinement on forests of octrees.

SIAM J. Sci. Comput., 33 no. 3 (2011), pages 1103-1133.


Triangulation

- Partitioning is cheap and simple:




- Then: take *p4est* refinement information
- Recreate rich *deal.II* Triangulation only for local cells (stores coordinates, connectivity, faces, materials, ...)
- How? recursive queries to *p4est*
- Also create ghost layer (one layer of cells around own ones)


Example (color by CPU ID)




=



&



&



What's needed?

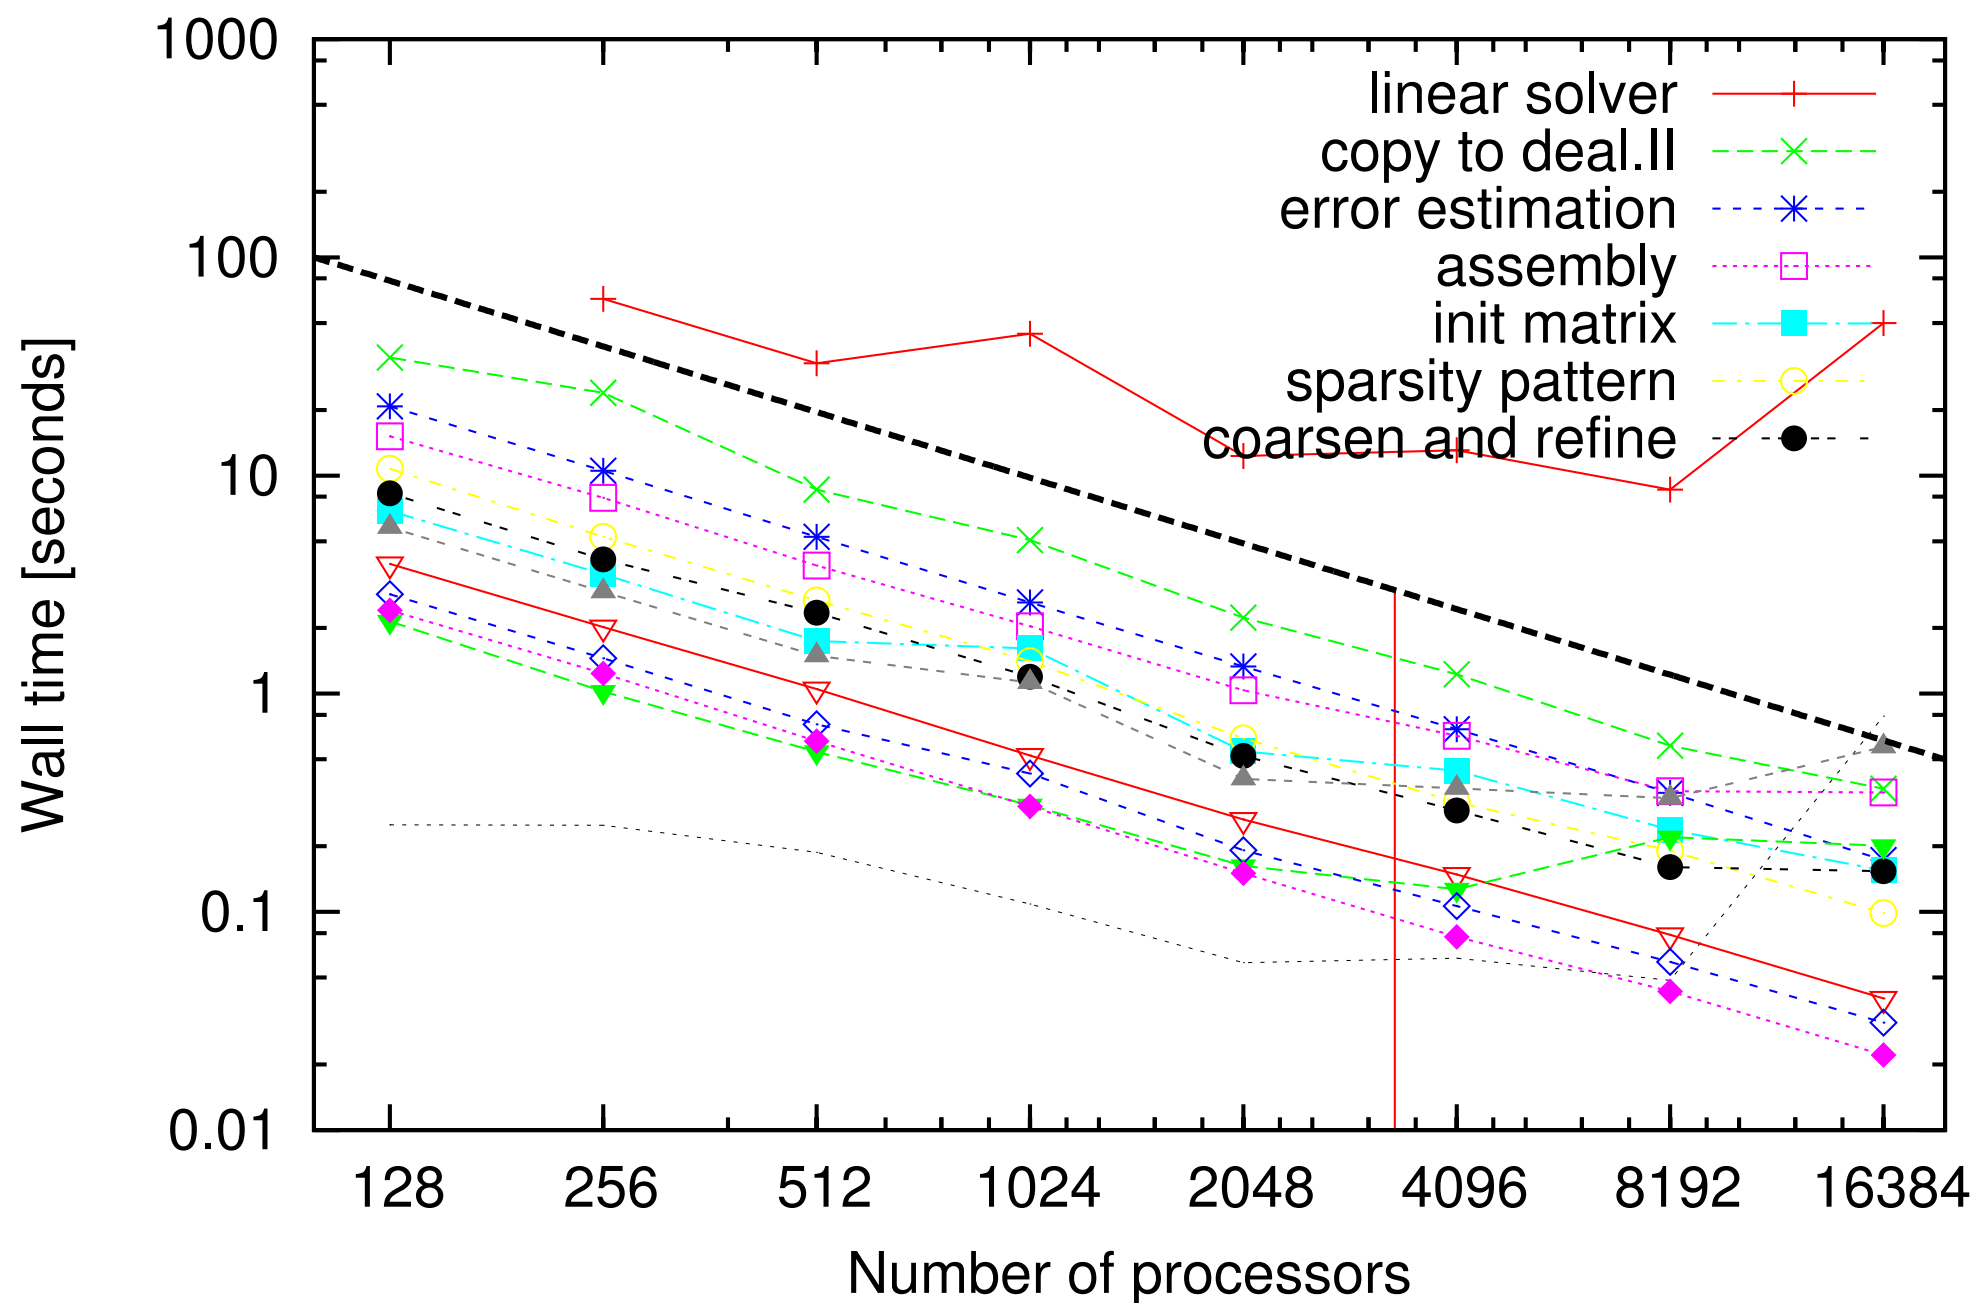
	serial mesh	dynamic parallel mesh	static parallel mesh
name	Triangulation	parallel::distributed ::Triangulation	(just an idea)
duplicate	everything	coarse mesh	nothing
partitioning	METIS	p4est: fast, scalable	offline, (PAR)METIS?
part. quality	good	okay	better?
hp?	yes	(planned)	yes?
geom. MG?	yes	in progress	?
Aniso. ref.?	yes	no	(offline only)
Periodicity	yes	in progress	?
Scalability	100 cores	16k+ cores	?

Solvers


- 🐾 Iterative solvers only need Mat-Vec products and scalar products
~> equivalent to serial code
- 🐾 Can use templated deal.II solvers like GMRES!
- 🐾 Better: use tuned parallel iterative solvers that hide/minimize communication
- 🐾 Preconditioners: more work, just operating on local blocks not enough

Strong Scaling: 2d Poisson

Wall clock times for problem of fixed size 335M



Memory Consumption




average and maximum memory consumption (VmPeak)
3D, weak scalability from 8 to 1000 processors with about 500.000
DoFs per processor (4 million up to 500 million total)



~> **Constant memory usage with increasing
CPUs & problem size**

Step 40

Strong Scaling (9.9M DoFs)



Weak Scaling (1.2M DoFs/Core)



Trilinos VS PETSc

What should I use?



- 🐾 Similar features and performance
- 🐾 Pro Trilinos: more development, some more features (automatic differentiation, . . .), cooperation with deal.II
- 🐾 Pro PETSc: stable, easier to compile on older clusters
- 🐾 But: being flexible would be better! – “why not both?”
 - 🐾 you can! Example: new step-40
 - 🐾 can switch at compile time
 - 🐾 need `#ifdef` in a few places (different solver parameters TrilinosML vs BoomerAMG)
 - 🐾 some limitations, somewhat work in progress

Trilinos VS PETSc

```
#include <deal.II/lac/generic_linear_algebra.h>
#define USE_PETSC_LA // uncomment this to run with Trilinos

namespace LA
{
#ifdef USE_PETSC_LA
    using namespace dealii::LinearAlgebraPETSc;
#else
    using namespace dealii::LinearAlgebraTrilinos;
#endif
}

// ...
LA::MPI::SparseMatrix system_matrix;
LA::MPI::Vector solution;

// ...
LA::SolverCG solver(solver_control, mpi_communicator);
LA::MPI::PreconditionAMG preconditioner;

LA::MPI::PreconditionAMG::AdditionalData data;

#ifdef USE_PETSC_LA
    data.symmetric_operator = true;
#else
    // trilinos defaults are good
#endif
    preconditioner.initialize(system_matrix, data);

// ...
```